

wine-quality-analysis-using-regression

September 26, 2023

```
[1]: # This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python Docker image: https://github.com/kaggle/
    ↪ docker-python
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list
    ↪ all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that gets
    ↪ preserved as output when you create a version using "Save & Run All"
# You can also write temporary files to /kaggle/temp/, but they won't be saved
    ↪ outside of the current session
```

```
[2]: import pandas as pd
import matplotlib.pyplot as plt
```

```
[3]: df = pd.read_csv('WineQT.csv')
df
```

```
[3]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.700	0.00	1.9	0.076	
1	7.8	0.880	0.00	2.6	0.098	
2	7.8	0.760	0.04	2.3	0.092	
3	11.2	0.280	0.56	1.9	0.075	
4	7.4	0.700	0.00	1.9	0.076	
...	
1138	6.3	0.510	0.13	2.3	0.076	
1139	6.8	0.620	0.08	1.9	0.068	

1140	6.2	0.600	0.08	2.0	0.090
1141	5.9	0.550	0.10	2.2	0.062
1142	5.9	0.645	0.12	2.0	0.075

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates \
0	11.0	34.0	0.99780	3.51	0.56
1	25.0	67.0	0.99680	3.20	0.68
2	15.0	54.0	0.99700	3.26	0.65
3	17.0	60.0	0.99800	3.16	0.58
4	11.0	34.0	0.99780	3.51	0.56
...
1138	29.0	40.0	0.99574	3.42	0.75
1139	28.0	38.0	0.99651	3.42	0.82
1140	32.0	44.0	0.99490	3.45	0.58
1141	39.0	51.0	0.99512	3.52	0.76
1142	32.0	44.0	0.99547	3.57	0.71

	alcohol	quality	Id
0	9.4	5	0
1	9.8	5	1
2	9.8	5	2
3	9.8	6	3
4	9.4	5	4
...
1138	11.0	6	1592
1139	9.5	6	1593
1140	10.5	5	1594
1141	11.2	6	1595
1142	10.2	5	1597

[1143 rows x 13 columns]

```
[4]: df.describe()
```

```
[4]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar \
count	1143.000000	1143.000000	1143.000000	1143.000000
mean	8.311111	0.531339	0.268364	2.532152
std	1.747595	0.179633	0.196686	1.355917
min	4.600000	0.120000	0.000000	0.900000
25%	7.100000	0.392500	0.090000	1.900000
50%	7.900000	0.520000	0.250000	2.200000
75%	9.100000	0.640000	0.420000	2.600000
max	15.900000	1.580000	1.000000	15.500000

	chlorides	free sulfur dioxide	total sulfur dioxide	density \
count	1143.000000	1143.000000	1143.000000	1143.000000
mean	0.086933	15.615486	45.914698	0.996730

std	0.047267	10.250486	32.782130	0.001925
min	0.012000	1.000000	6.000000	0.990070
25%	0.070000	7.000000	21.000000	0.995570
50%	0.079000	13.000000	37.000000	0.996680
75%	0.090000	21.000000	61.000000	0.997845
max	0.611000	68.000000	289.000000	1.003690

	pH	sulphates	alcohol	quality	Id
count	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000
mean	3.311015	0.657708	10.442111	5.657043	804.969379
std	0.156664	0.170399	1.082196	0.805824	463.997116
min	2.740000	0.330000	8.400000	3.000000	0.000000
25%	3.205000	0.550000	9.500000	5.000000	411.000000
50%	3.310000	0.620000	10.200000	6.000000	794.000000
75%	3.400000	0.730000	11.100000	6.000000	1209.500000
max	4.010000	2.000000	14.900000	8.000000	1597.000000

```
[5]: df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1143 entries, 0 to 1142
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1143 non-null   float64
1   volatile acidity       1143 non-null   float64
2   citric acid            1143 non-null   float64
3   residual sugar         1143 non-null   float64
4   chlorides              1143 non-null   float64
5   free sulfur dioxide    1143 non-null   float64
6   total sulfur dioxide   1143 non-null   float64
7   density                1143 non-null   float64
8   pH                    1143 non-null   float64
9   sulphates              1143 non-null   float64
10  alcohol                1143 non-null   float64
11  quality                1143 non-null   int64
12  Id                     1143 non-null   int64
dtypes: float64(11), int64(2)
memory usage: 116.2 KB
```

0.0.1 no null values


0.0.2 visulize for null values

```
[6]: import seaborn as sns
sns.heatmap(df.isnull())
```

```
[6]: <AxesSubplot: >
```



```
[7]: # visualize count of quality
sns.countplot(data = df, x = df['quality'], color= sns.color_palette()[9])
plt.show()
```



```
[8]: #input data
x = df.iloc[:, :11]
x.head()
```

```
[8]: fixed acidity  volatile acidity  citric acid  residual sugar  chlorides  \
0          7.4           0.70           0.00           1.9           0.076
1          7.8           0.88           0.00           2.6           0.098
2          7.8           0.76           0.04           2.3           0.092
3         11.2           0.28           0.56           1.9           0.075
4          7.4           0.70           0.00           1.9           0.076

    free sulfur dioxide  total sulfur dioxide  density    pH  sulphates  \
0             11.0           34.0    0.9978  3.51      0.56
1             25.0           67.0    0.9968  3.20      0.68
2             15.0           54.0    0.9970  3.26      0.65
3             17.0           60.0    0.9980  3.16      0.58
4             11.0           34.0    0.9978  3.51      0.56

    alcohol
0       9.4
1       9.8
```

```

2      9.8
3      9.8
4      9.4

```

```

[9]: #output data
y = df.iloc[:,11:12]
y.head()

```

```

[9]:    quality
0      5
1      5
2      5
3      6
4      5

```

0.1 split data into train and test

```

[10]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size= .
↪25,random_state=12)

```

```

[11]: x_train.head()

```

```

[11]:    fixed acidity  volatile acidity  citric acid  residual sugar  chlorides \
1038          7.9           0.200         0.35         1.7         0.054
638          7.4           0.635         0.10         2.4         0.080
765          8.5           0.340         0.40         4.7         0.055
81           9.4           0.400         0.31         2.2         0.090
489          9.2           0.920         0.24         2.6         0.087

    free sulfur dioxide  total sulfur dioxide  density    pH  sulphates \
1038                7.0                15.0  0.99458  3.32        0.80
638                16.0                33.0  0.99736  3.58        0.69
765                 3.0                 9.0  0.99738  3.38        0.66
81                13.0                62.0  0.99660  3.07        0.63
489                12.0                93.0  0.99980  3.48        0.54

    alcohol
1038    11.9
638    10.8
765    11.6
81     10.5
489     9.8

```

0.2 scale the data

```
[12]: from sklearn.preprocessing import StandardScaler
      sc = StandardScaler()
      x_train = sc.fit_transform(x_train)
      x_test = sc.fit_transform(x_test)
```

```
[13]: x_train
```

```
[13]: array([[ -0.23190535, -1.84245871,  0.45326153, ...,  0.02751629,
          0.83401434,  1.42166275],
        [ -0.52374511,  0.55407797, -0.84588607, ...,  1.70270196,
          0.18723471,  0.37290974],
        [  0.11830237, -1.07115955,  0.71309105, ...,  0.4140976 ,
          0.01084027,  1.1356392 ],
        ...,
        [ -0.34864125,  0.85708836, -0.22229522, ..., -0.03691393,
         -1.0475264 , -0.86652564],
        [ -0.99068874,  1.10500594, -1.2616133 , ...,  0.99396956,
         -0.75353566, -0.48516091],
        [  0.29340624, -1.62208752,  1.44061371, ...,  0.47852782,
         -0.69473751, -1.15254919]])
```

```
[14]: x_test
```

```
[14]: array([[ 0.83963515,  0.04571776,  1.37174921, ..., -0.66865691,
          0.6058255 ,  0.90524944],
        [ 0.08005559, -1.97474479,  1.08285357, ...,  0.20713709,
          0.6058255 ,  1.07632962],
        [ -0.19122283, -1.57065228,  0.69765939, ...,  0.51992066,
         -0.32948404,  1.16186971],
        ...,
        [ -0.84229102,  0.96935778, -1.37275933, ...,  1.02037438,
          0.37199811, -0.29231182],
        [ 0.08005559, -0.18519225, -0.84311733, ..., -0.79377034,
         -0.73868197, -0.97663254],
        [ -0.24547851, -1.10883227,  0.60136085, ..., -0.04308976,
          0.31354127,  0.3920089 ]])
```

1 LinearRegression

```
[15]: from sklearn.linear_model import LinearRegression
      lr = LinearRegression()
      lr.fit(x_train, y_train)
```

```
[15]: LinearRegression()
```

```
[16]: #Predicting the test set result  
y_pred = lr.predict(x_test)
```

```
[17]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score  
import numpy as np  
# Calculate evaluation metrics  
mae = mean_absolute_error(y_test, y_pred)  
mse = mean_squared_error(y_test, y_pred)  
rmse = np.sqrt(mse)  
r2 = r2_score(y_test, y_pred)  
  
print("Mean Absolute Error:", mae)  
print("Mean Squared Error:", mse)  
print("Root Mean Squared Error:", rmse)  
print("R-squared:", r2)
```

```
Mean Absolute Error: 0.5143692288772529  
Mean Squared Error: 0.4466318515999171  
Root Mean Squared Error: 0.6683052084189657  
R-squared: 0.28434612651882896
```

2 Support Vector Regressor

```
[18]: from sklearn.svm import SVR  
svr = SVR(kernel='rbf')  
svr.fit(x_train,y_train)
```

```
C:\Users\shivs\anaconda3\Lib\site-packages\sklearn\utils\validation.py:1184:  
DataConversionWarning: A column-vector y was passed when a 1d array was  
expected. Please change the shape of y to (n_samples, ), for example using  
ravel().  
y = column_or_1d(y, warn=True)
```

```
[18]: SVR()
```

```
[19]: y_pred_svr = svr.predict(x_test)  
y_pred_svr
```

```
[19]: array([6.07837003, 6.21826791, 6.06904514, 5.51922936, 5.49701865,  
        5.0464841 , 5.85929132, 4.98278073, 5.1664285 , 6.58155529,  
        5.46525041, 5.1854067 , 5.01439475, 4.98823248, 5.09041047,  
        4.9491099 , 5.89835232, 5.84271973, 5.15769836, 6.30728981,  
        5.25079339, 5.11742525, 4.96345776, 5.55225868, 5.57779887,  
        5.71717254, 5.71702483, 5.02673959, 5.62610822, 4.92411501,  
        5.57094214, 5.63706285, 5.18218594, 5.18288003, 5.64892325,  
        5.91944208, 5.09200407, 5.26596778, 5.5188043 , 5.66189109,  
        5.60674896, 5.66189109, 4.95561156, 5.7344404 , 5.8977396 ,  
        4.91987445, 5.93325773, 5.8172807 , 5.45154511, 6.25583049,
```


5.76743799, 6.21284083, 5.34438289, 5.70556665, 6.56072001,
6.57440205, 5.90523183, 5.54476259, 6.31490625, 4.95384797,
5.56230762, 6.60520854, 5.97264289, 5.65388327, 5.92529893,
5.16445224, 4.92008195, 6.10151423, 4.98698548, 5.90192542,
5.12279306, 5.85675985, 5.57175204, 5.8597259 , 5.91216414,
4.95126465, 6.24820242, 4.88394445, 5.85000763, 5.67003837,
5.20885606, 5.8919132 , 5.73204691, 6.10255564, 5.64319119,
4.73133567, 5.3838868 , 5.39135041, 5.46405134, 6.22181804,
5.31143196, 5.83319738, 5.12380685, 5.09155609, 5.91865445,
4.92861213, 6.08852023, 6.23200664, 5.58996514, 6.458952 ,
5.73038932, 5.32125964, 5.74097941, 5.70556665, 4.95384797,
5.33495576, 5.70172189, 5.37002236, 6.68142775, 6.7618333 ,
5.30103648, 5.13533283, 6.37027081, 5.9229749 , 5.33970578,
5.3266891 , 5.66510044, 6.20613738, 6.68880678, 6.479259 ,
6.11621348, 5.2303019 , 7.12766649, 4.85362213, 5.65693876,
5.42227329, 6.11363992, 6.22387461, 5.24589517, 5.51960768,
5.82847846, 5.62704069, 5.73131164, 5.70417154, 5.18720504,
5.18218594, 5.72884243, 5.16873236, 5.38439838, 6.9409716 ,
5.21592012, 6.97578144, 5.46438735, 6.03245621, 4.969765 ,
6.92325775, 6.21707597, 5.62374367, 5.62610822, 4.68327915,
5.46833266, 5.0461287 , 5.01505039, 5.25786767, 5.29579624,
5.27912451, 5.37498789, 4.86426885, 4.9093795 , 5.07673976,
6.56235532, 4.96168843, 6.64474326, 5.52727373, 5.67674338,
6.83946077, 6.41378527, 5.85830233, 5.25718578, 5.95515385,
5.02324894, 5.75584941, 5.50844586, 5.49749899, 4.98934389,
6.08694428, 6.02759942, 5.16358006, 5.32478633, 4.99155161,
5.82453786, 5.36393021, 6.01345298, 5.21445665, 5.34511376,
5.27430361, 6.5723734 , 5.12060014, 5.46020247, 6.13694716,
5.43988205, 5.07688471, 5.99386761, 5.075961 , 6.08327897,
6.10380093, 5.17646139, 6.16585042, 5.11728571, 5.08219642,
5.2643713 , 5.13979551, 5.37518733, 5.21137363, 5.87138566,
5.60206626, 5.01108803, 5.49026155, 6.40673387, 5.52558597,
5.50424672, 5.10389012, 5.91709832, 5.84783973, 6.40804598,
5.04673667, 5.22347667, 5.17272174, 5.64563315, 6.2992712 ,
5.19577953, 6.84031999, 5.02871918, 6.15889257, 5.84261355,
5.71978398, 5.04857979, 5.98444988, 6.08571055, 5.59283342,
6.86239995, 5.09639639, 5.65334324, 5.30666866, 5.17947265,
6.46929686, 5.32537935, 4.88347585, 5.91433563, 5.92286152,
4.79344064, 4.83731738, 5.24264032, 5.53695642, 6.01565189,
4.95773032, 6.2762183 , 5.43988205, 5.56615368, 5.4402394 ,
5.54925322, 5.05877033, 6.42155046, 5.91579167, 6.5802308 ,
5.1796325 , 6.41881052, 6.08795147, 6.82741404, 5.37239931,
5.01757471, 5.49854768, 4.91436746, 6.65446805, 4.91056632,
5.10785202, 6.46357597, 5.79205622, 6.35300457, 6.10443234,
5.27957812, 5.65766101, 5.78030912, 6.1301421 , 6.40673387,
6.6287756 , 6.35449889, 5.43556546, 5.74692251, 5.91709832,
5.81100503, 6.2555767 , 6.14029902, 5.47493124, 5.30542261,

6.15354495])

```
[20]: # Calculate evaluation metrics
mae_svr = mean_absolute_error(y_test, y_pred_svr)
mse_svr = mean_squared_error(y_test, y_pred_svr)
rmse_svr = np.sqrt(mse_svr)
r2_svr = r2_score(y_test, y_pred_svr)

print("Mean Absolute Error:", mae_svr)
print("Mean Squared Error:", mse_svr)
print("Root Mean Squared Error:", rmse_svr)
print("R-squared:", r2_svr)
```

Mean Absolute Error: 0.47655708362025173
Mean Squared Error: 0.41097827040612117
Root Mean Squared Error: 0.6410758694617363
R-squared: 0.3414751095804127

3 Decision Tree Regressor

```
[21]: from sklearn.tree import DecisionTreeRegressor
dtr = DecisionTreeRegressor(random_state=12)
dtr.fit(x_train, y_train)
```

[21]: DecisionTreeRegressor(random_state=12)

```
[22]: y_pred_dtr = dtr.predict(x_test)
y_pred_dtr
```

```
[22]: array([6., 5., 6., 5., 5., 5., 6., 5., 6., 5., 5., 4., 5., 4., 5., 5., 6.,
        8., 5., 7., 5., 5., 5., 5., 6., 5., 6., 5., 5., 5., 6., 7., 5., 5.,
        6., 5., 6., 6., 5., 5., 6., 5., 5., 5., 6., 6., 5., 7., 6., 6., 5.,
        6., 6., 6., 7., 7., 6., 5., 6., 5., 5., 7., 4., 6., 6., 5., 5., 6.,
        4., 6., 5., 5., 6., 6., 6., 5., 8., 5., 7., 5., 5., 6., 5., 5., 6.,
        5., 5., 5., 5., 6., 5., 6., 5., 5., 4., 5., 6., 6., 6., 6., 5., 5.,
        7., 6., 5., 6., 7., 5., 7., 7., 5., 5., 6., 6., 5., 6., 6., 6., 7.,
        6., 7., 5., 6., 6., 6., 4., 7., 6., 5., 6., 6., 6., 6., 5., 5., 5.,
        5., 5., 5., 7., 6., 7., 5., 6., 5., 8., 6., 5., 5., 5., 6., 5., 5.,
        6., 5., 5., 5., 5., 5., 5., 8., 5., 6., 6., 6., 7., 6., 5., 5., 6.,
        6., 6., 5., 5., 5., 6., 6., 4., 5., 5., 6., 5., 6., 5., 5., 5., 6.,
        5., 6., 6., 5., 5., 6., 5., 6., 7., 5., 6., 5., 5., 5., 4., 6., 5.,
        4., 5., 5., 5., 6., 5., 6., 6., 5., 5., 7., 5., 6., 5., 5., 7., 5.,
        6., 5., 5., 6., 6., 5., 6., 6., 7., 7., 5., 6., 5., 5., 5., 4., 3.,
        7., 6., 5., 5., 5., 6., 6., 5., 7., 5., 6., 6., 6., 5., 7., 6., 7.,
        6., 5., 5., 7., 6., 5., 6., 5., 6., 5., 6., 6., 7., 7., 7., 5., 7.,
        6., 6., 6., 8., 6., 6., 6., 5., 7., 6., 6., 6., 6., 7.]])
```

```
[23]: # Calculate evaluation metrics
mae_dtr = mean_absolute_error(y_test, y_pred_dtr)
mse_dtr = mean_squared_error(y_test, y_pred_dtr)
rmse_dtr = np.sqrt(mse_dtr)
r2_dtr = r2_score(y_test, y_pred_dtr)

print("Mean Absolute Error:", mae_dtr)
print("Mean Squared Error:", mse_dtr)
print("Root Mean Squared Error:", rmse_dtr)
print("R-squared:", r2_dtr)
```

Mean Absolute Error: 0.5559440559440559
Mean Squared Error: 0.7657342657342657
Root Mean Squared Error: 0.8750624353349112
R-squared: -0.22696285848613074

In summary, lower values for MAE, MSE, and RMSE are desirable, as they indicate better predictive performance. A higher R^2 value indicates a better fit of the model to the data, with values closer to 1 being better. Algorithm 2 seems to perform better than Algorithm 1 and Algorithm 3 based on these metrics.

4 Perform k-fold cross-validation

1. Linear Regression

```
[24]: from sklearn.model_selection import cross_val_score
lr_k = LinearRegression()
k = 5
mse_scores = -cross_val_score(lr_k, x, y, cv=k, scoring='neg_mean_squared_error')
```

we use - here because of we use scoring='neg_mean_squared_error' this

```
[25]: mse_scores
```

```
[25]: array([0.44290968, 0.46117682, 0.40050413, 0.40189841, 0.4247529 ])
```

2. Support Vector Regressor

```
[26]: from sklearn.model_selection import GridSearchCV, KFold

# Define the parameter grid
param_grid = {'C': [0.01, 0.1, 1, 10], 'gamma': [0.0001, 0.0005, 0.001, 0.005, 0.01, 0.05]}
```

```
[27]: # Create an SVR model
svr_k = SVR()

# Initialize KFold with k splits
k = 5
```

```
kf = KFold(n_splits=k, shuffle=True, random_state=12)
```

```
[28]: # Perform grid search with cross-validation
grid_search = GridSearchCV(svr_k, param_grid, cv=kf,
    ↳scoring='neg_mean_squared_error', n_jobs=-1)
grid_search.fit(x, y.values.ravel())
```

```
[28]: GridSearchCV(cv=KFold(n_splits=5, random_state=12, shuffle=True),
    estimator=SVR(), n_jobs=-1,
    param_grid={'C': [0.01, 0.1, 1, 10],
    'gamma': [0.0001, 0.0005, 0.001, 0.005, 0.01, 0.05]},
    scoring='neg_mean_squared_error')
```

- GridSearchCV is a technique used for hyperparameter tuning in machine learning. It's employed to find the best combination of hyperparameters for a machine learning model, leading to improved model performance.
- cv=kf: This parameter specifies the cross-validation strategy to use during the grid search
- n_jobs=-1: This parameter specifies the number of CPU cores to use for parallel computation. Setting it to -1 means to use all available CPU cores, which can significantly speed up the grid search if you have a multi-core processor.

```
[29]: # Get the best hyperparameters
grid_search.best_params_
```

```
[29]: {'C': 10, 'gamma': 0.005}
```

```
[30]: # Get the best estimator (SVR model with best hyperparameters)
best_model = grid_search.best_estimator_
```

```
[31]: # Calculate the average Mean Squared Error across all folds using the best model
mse_scores = -cross_val_score(best_model, x, y.values.ravel(), cv=k,
    ↳scoring='neg_mean_squared_error')
```

```
[32]: mse_scores
```

```
[32]: array([0.46114719, 0.54387765, 0.3905256 , 0.51144096, 0.49882219])
```

- now get best score 0.390526. This is the lowest MSE value, and therefore, it represents the best performance on the third fold of cross-validation.

5 Improvement

```
[33]: # train test split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=
    ↳20, random_state=12)
```

```
[34]: # Define the parameter grid
param_grid_1 = {'C': [0.01, 0.1, 1, 10], 'gamma': [0.0001, 0.0005, 0.001, 0.005, 0.01, 0.05]}
```

```
[35]: # Create an SVR model
svr_k = SVR()

# Initialize KFold with k splits
k = 5
kf = KFold(n_splits=k, shuffle=True, random_state=12)
```

```
[36]: # Perform grid search with cross-validation
grid_search = GridSearchCV(svr_k, param_grid_1, cv=kf,
    scoring='neg_mean_squared_error', n_jobs=-1)
grid_search.fit(x_train, y_train.values.ravel())
```

```
[36]: GridSearchCV(cv=KFold(n_splits=5, random_state=12, shuffle=True),
    estimator=SVR(), n_jobs=-1,
    param_grid={'C': [0.01, 0.1, 1, 10],
    'gamma': [0.0001, 0.0005, 0.001, 0.005, 0.01, 0.05]},
    scoring='neg_mean_squared_error')
```

```
[37]: # Get the best estimator (SVR model with best hyperparameters)
best_model = grid_search.best_estimator_
```

```
[38]: # Calculate the average Mean Squared Error across all folds using the best model
mse_scores = -cross_val_score(best_model, x_train, y_train.values.ravel(),
    cv=kf, scoring='neg_mean_squared_error')
mse_scores.mean()
```

```
[38]: 0.46535398740587636
```

- Average Mean Squared Error (Cross-Validation): 0.46535398740587597

```
[39]: # Calculate Mean Squared Error on the test set using the best model
y_pred = best_model.predict(x_test)
mse_test = mean_squared_error(y_test, y_pred)
mse_test
```

```
[39]: 0.49675628780242
```

- Mean Squared Error on Test Set: 0.49675628780242076